

# Uncertainty-Aware Bayesian Mineral Prospectivity Mapping under Spatial Separation: Heterogeneous Predictor Relationships across the Central African Copperbelt

## Abstract

Mineral prospectivity mapping integrates geological, geochemical, and geophysical evidence to identify locations whose mapped characteristics are associated with known deposits. We evaluate a partially pooled hierarchical Bayesian logistic model across four geographically contiguous held-out sectors of the Central African Copperbelt. The 1,872-cell modeling frame consists of a regular 5 km  $\times$  5 km grid containing 138 documented sediment-hosted stratiform copper-cobalt occurrences, with predictors for distance to mapped faults, distance to lithological contacts, Bouguer gravity, and retained host-lithology classes. Intercepts and the four distance coefficients vary by Daly tectonic domain; Bouguer gravity and retained host-rock effects are global. Posterior parameter uncertainty is quantified; cell-level predictive intervals and Bernoulli outcome draws are not reported. Each held-out cell receives the posterior mean of its draw-wise conditional probabilities  $\bar{p}_i = \frac{1}{S} \sum_s p_i^{(s)}$ . No classification threshold (e.g., 0.5) is applied to generate predictions or compute discrimination metrics; performance is evaluated across the continuous probability distribution. Predictor contributions to the decomposed out-of-fold (OOF) linear predictor vary across folds, and across four distance-representation specifications mean ROC-AUC ranges from 0.671 to 0.686 while mean PR-AUC ranges from 0.107 to 0.145. Distance associations recur across evaluated representations, whereas quadratic curvature is representation-dependent and population-level  $D^*$  is not identifiable. Bootstrap intervals quantify uncertainty in performance metrics, not individual cell probabilities. The results demonstrate geographic variation in fitted associations and held-out discrimination, but do not establish distinct causal geological processes or a universal distance law.

**Keywords:** mineral prospectivity mapping; Bayesian hierarchical model; spatial cross-validation; spatial heterogeneity; Central African Copperbelt; Daly tectonic domains; parameter uncertainty

## 1 Introduction

Mineral prospectivity mapping integrates geological, geochemical, and geophysical evidence to identify locations whose mapped characteristics are associated with known deposits. Published approaches range from knowledge-driven evidence integration (Bonham-Carter, 1994; Carranza, 2009) to logistic regression, weights of evidence (Bonham-Carter et al., 1989), Bayesian networks, and machine-learning classifiers (Porwal et al., 2006; Carranza and Laborte, 2015; Ford, 2019; Yousefi et al., 2021; Matthieu et al., 2026); they differ in assumptions, data requirements, and treatment of incomplete deposit inventories. In the Central African Copperbelt, stratigraphy, structures, basin evolution, and ore assemblages vary regionally, while the timing and detailed genesis of mineralization remain debated (Cailteux et al., 2005; Selley et al., 2005; Hitzman et al., 2012; Key et al.,

2004; Ndonfack et al., 2025). The predictors used here are spatial proxies within that geological system, not direct measurements of ore-forming processes.

Spatial structure also affects evaluation. When nearby observations share covariates, occurrence history, or sampling processes, random cross-validation can place dependent observations in both training and test sets and underestimate error for a new-region prediction task (Brenning, 2012; Roberts et al., 2017). The size and direction of this bias depend on the prediction target and blocking design; spatial blocks can also create predictor extrapolation (Valavi et al., 2019). This study therefore evaluates transfer among four contiguous along-belt sectors rather than treating random cell splitting as a proxy for geographic transfer.

The Bayesian formulation represents posterior uncertainty in shared parameters, hierarchical scales, and domain effects. Draw-wise held-out conditional probabilities are averaged to produce stored OOF scores, but the frozen OOF artifact does not retain cell-level probability intervals. Future Bernoulli outcomes are not sampled. Uncertainty in evidence layers, model parameters, conditional probabilities, and outcomes are distinct quantities in prospectivity work (Lisitsin et al., 2014; Parsa and Carranza, 2021; Zuo et al., 2021).

The Central African Copperbelt provides a setting in which to examine regional transfer. Daly’s domain labels are used as a geological stratification for model terms and secondary summaries (Daly et al., 2025); the model estimates associations indexed by those labels but does not validate the tectonic framework itself. Spatially varying coefficient models offer a statistical representation of changing relationships without making that variation causal (Gelfand et al., 2003; Fotheringham et al., 2002).

This study contributes: (i) a partially pooled hierarchical Bayesian mineral prospectivity framework evaluated under along-belt spatial separation to test geographic transferability without spatial leakage; and (ii) a representation-sensitivity design that explicitly separates robust, recurring proximity associations from representation-dependent polynomial curvature and  $D^*$  stationary-point diagnostics.

The objective is to estimate multivariate geological and geophysical predictor associations under a hierarchical Bayesian model and evaluate how fitted relationships and discrimination vary across four spatially separated Copperbelt sectors. We ask how fault distance, lithological-contact distance, Bouguer gravity, and host lithology jointly relate to the recorded occurrence labels; how model-component contributions differ across held-out sectors; how OOF discrimination varies; and whether distance-response conclusions persist under alternative mathematical representations.

## 2 Study area, spatial data, and evidence inventory

### 2.1 Study area and grid definition

The study area encompasses the Central African Copperbelt tract represented in the project spatial modelling grid, spanning the border region of northern Zambia and the southern Democratic Republic of the Congo. The spatial frame is defined on a regular  $5\text{ km} \times 5\text{ km}$  grid projected in WGS 84 / UTM Zone 35S (EPSG:32735). The grid extends from Easting 282,675 m to 715,465 m and Northing 8,380,251 m to 8,859,512 m, covering an along-belt tract of approximately 500 km in strike length.

The observational unit is a grid cell identified by its centroid coordinates  $(x_i, y_i)$  and a binary

occurrence indicator  $Y_i \in \{0, 1\}$ . Cells containing at least one documented sediment-hosted stratiform copper-cobalt deposit or significant occurrence are labeled  $Y_i = 1$  ( $n = 138$ ). Cells without a recorded occurrence in the compiled inventory are labeled  $Y_i = 0$  ( $n = 1,734$ ). In mineral exploration, a zero denotes the absence of a recorded occurrence in the compiled database, not confirmed barren ground. After complete-case filtering across all continuous and categorical predictors, the modelling frame comprises exactly 1,872 cells.

## 2.2 Data sources and feature extraction

Predictor variables were compiled from regional geological, structural, geophysical, and tectonic layers:

1. **Distance to mapped faults ( $d_f$ ):** Regional fault lineaments and major structural boundaries were compiled from regional structural maps of the Katangan basin (**CACB.Structures**). For each grid cell,  $d_f$  was computed as the minimum Euclidean distance (in meters) from the cell centroid to the nearest mapped fault trace.
2. **Distance to lithological contacts ( $d_l$ ):** Stratigraphic contacts between Neoproterozoic Katangan sedimentary formations and underlying basement or adjacent units were extracted from regional geological line work (**all\_lithology\_lines**). The distance  $d_l$  was computed as the minimum Euclidean distance from the cell centroid to the nearest contact trace.
3. **Bouguer gravity ( $g$ ):** Regional Bouguer gravity anomaly values were extracted from regional terrestrial-satellite gravity grids (WGM2012 / Bureau Gravimétrique International), sampled at each cell centroid in milligals (mGal), reflecting crustal density variations and basement architecture.
4. **Host-rock lithology ( $\mathbf{x}_{\text{rock}}$ ):** Surficial lithological parent material classes were derived from the African Surficial Lithology dataset compiled by the U.S. Geological Survey and RCMRD (Sayre et al., 2014), which classifies bedrock geology and surficial deposits into 19 regional units. The dominant class at each cell centroid was assigned.
5. **Tectonic domain stratification ( $d(i)$ ):** Regional tectonic domain polygons were digitized from the recent tectonic framework of the Katangan basin by Daly et al. (2025). Each cell was assigned to one of six Daly domains via spatial point-in-polygon intersection: Northern Red Beds domain 3a (NRB\_3a), Northern Red Beds domain 3b (NRB\_3b), Central Ridge Zone (CRZ), Northern Kundelungu Basin (NKB), Southern Red Beds (SRB), and Mwashya–Mulungushi Sub-basin (MMSB).

## 2.3 Domain-level positive-observation structure and partial-pooling qualification

Documented deposit occurrences in the modelling frame are concentrated exclusively in two of the six tectonic domains: 107 positive cells occur in NRB\_3a (out of 1,127 total cells) and 31 occur in NRB\_3b (out of 550 total cells). The remaining four domains contain zero recorded positive observations in this compiled frame: CRZ ( $n = 154$ ), NKB ( $n = 24$ ), MMSB ( $n = 16$ ), and SRB ( $n = 1$ ).

Methodologically, this empirical structure has specific implications for the hierarchical Bayesian model:

- The formulation remains a fully partially pooled hierarchical model across all six domains, with domain-varying parameters linked through shared population hyperparameters  $\mu_p$  and  $\sigma_p$ .
- However, because CRZ, NKB, SRB, and MMSB contain no positive observations, their domain-specific coefficients are not informed by within-domain positive-versus-background contrast in the empirical likelihood.
- Instead, parameter estimates in these four single-class domains are governed almost entirely by hierarchical shrinkage toward the population mean  $\mu_p$  and by the shared background likelihood contributions.
- Consequently, domain-specific coefficient interpretations for CRZ, NKB, SRB, and MMSB must be treated with substantial caution, and standard binary classification metrics such as ROC-AUC are undefined within these single-class partitions.

## 2.4 Complete preprocessing chain

The complete preprocessing and spatial prediction workflow is shown in Figure 1. All fold-specific transformations—standardization of continuous predictors, construction of quadratic terms, and host-lithology support filtering—are fitted using training cells only and applied unchanged to the held-out cells, enforcing zero data leakage between spatial partitions. Formally, each continuous predictor is standardized as  $z_i = (X_i - \mu_{\text{train}})/\sigma_{\text{train}}$ , where the training-fold statistics are frozen before any held-out transformation. Quadratic terms  $z_{f,i}^2$  and  $z_{l,i}^2$  are constructed in  $z$ -space after standardization.

## 2.5 Distinction between validation folds and tectonic domains

It is essential to distinguish the two spatial stratifications used in this study:

- **Along-belt validation folds (Folds 1–4):** A purely geometric, 1D along-strike partition designed to evaluate out-of-fold spatial predictive transferability and prevent spatial auto-correlation leakage.
- **Daly tectonic domains (6 domains):** A geological and structural stratification based on Katangan basin evolution (Daly et al., 2025), used as the grouping variable for hierarchical partial pooling of model parameters.

The validation is neither Leave-One-Daly-Domain-Out nor a Fold  $\times$  Domain cross-validation. The model is trained on three spatial folds and evaluated on the held-out spatial fold; domain-level performance metrics are computed post hoc as descriptive strata of the spatial predictions.




Figure 1: Complete preprocessing and spatial prediction workflow for the V11 hierarchical Bayesian prospectivity model. Spatial partitioning precedes fold-specific preprocessing; standardization and host-lithology support filtering are estimated from training cells only, while held-out cells are transformed using the frozen training-fold operations.




Figure 2: Four along-belt contiguous spatial validation folds across the Central African Copperbelt modeling grid. Cells are partitioned into four balanced blocks of 468 cells along the first principal component of coordinates, enforcing spatial separation between training and test sets.

### 3 Bayesian hierarchical methodology and spatial validation

#### 3.1 V11 model specification

For cell  $i$  in Daly domain  $d(i) \in \{1, \dots, 6\}$ , the V11 model is a partially pooled hierarchical Bayesian logistic regression (Gelman and Hill, 2007):

$$Y_i \sim \text{Bernoulli}(p_i), \quad \text{logit}(p_i) = \eta_i, \quad (1)$$

where the linear predictor  $\eta_i$  decomposes as:

$$\eta_i = \alpha_{d(i)} + \beta_{f,\text{lin},d(i)} z_{f,i} + \beta_{f,\text{sq},d(i)} z_{f,i}^2 + \beta_{l,\text{lin},d(i)} z_{l,i} + \beta_{l,\text{sq},d(i)} z_{l,i}^2 + \beta_g z_{g,i} + \mathbf{x}_{\text{rock},i}^\top \boldsymbol{\beta}_{\text{rock}}. \quad (2)$$

The model distinguishes:

- **Domain-varying parameters:** Domain intercepts  $\alpha_d$  and the four distance coefficients (fault linear  $\beta_{f,\text{lin},d}$ , fault quadratic  $\beta_{f,\text{sq},d}$ , lithological contact linear  $\beta_{l,\text{lin},d}$ , and lithological contact quadratic  $\beta_{l,\text{sq},d}$ ).
- **Global parameters:** The Bouguer gravity coefficient  $\beta_g$  and the host-lithology vector  $\boldsymbol{\beta}_{\text{rock}}$  are shared globally across all domains.

Domain-varying parameters use a non-centred hierarchical parameterization to ensure efficient MCMC exploration:

$$\beta_{p,d} = \mu_p + \sigma_p \cdot \text{offset}_{p,d}, \quad \text{for } p \in \{f_{\text{lin}}, f_{\text{sq}}, l_{\text{lin}}, l_{\text{sq}}\}, \quad (3)$$

and

$$\alpha_d = \mu_\alpha + \sigma_\alpha \cdot \text{offset}_{\alpha,d}. \quad (4)$$

Prior distributions are weakly informative:

$$\mu_p \sim \mathcal{N}(0, 1), \quad \sigma_p \sim \text{HalfNormal}(1), \quad (5)$$

$$\text{offset}_{p,d} \sim \mathcal{N}(0, 1), \quad \mu_\alpha \sim \mathcal{N}(\text{logit}(\hat{p}_{\text{train}}), 1), \quad (6)$$

$$\sigma_\alpha \sim \text{HalfNormal}(1), \quad \text{offset}_{\alpha,d} \sim \mathcal{N}(0, 1), \quad (7)$$

$$\beta_g \sim \mathcal{N}(0, 1), \quad \boldsymbol{\beta}_{\text{rock}} \sim \mathcal{N}(\mathbf{0}, \mathbf{I}), \quad (8)$$

where  $\hat{p}_{\text{train}} = \sum y_{\text{train}} / N_{\text{train}}$  anchors the intercept prior to the empirical training base rate.

#### 3.2 Continuous probability prediction and threshold-free evaluation

For each posterior draw  $s \in \{1, \dots, S\}$  and held-out cell  $i$ , the model computes draw-wise linear predictors and conditional probabilities:

$$\eta_i^{(s)} = \alpha_{d(i)}^{(s)} + \dots, \quad p_i^{(s)} = \text{logit}^{-1}(\eta_i^{(s)}) = \frac{1}{1 + \exp(-\eta_i^{(s)})}. \quad (9)$$

The reported out-of-fold prediction is the posterior mean conditional probability:

$$\bar{p}_i = \frac{1}{S} \sum_{s=1}^S p_i^{(s)}. \quad (10)$$

**No classification threshold (e.g., 0.5) was applied to generate the reported OOF predictions or calculate ROC-AUC/PR-AUC.** Predictive performance is evaluated across the continuous distribution of predicted probabilities using threshold-independent rank discrimination metrics (Lobo et al., 2008).

### 3.3 Uncertainty concepts

We maintain four distinct uncertainty concepts:

1. **Parameter posterior uncertainty:** Represented by MCMC posterior samples of coefficients and hyperparameters.
2. **Conditional probability uncertainty:** Represented by the distribution of draw-wise probabilities  $\{p_i^{(s)}\}_{s=1}^S$ . The reported OOF map presents the posterior mean  $\bar{p}_i$ ; cell-level credible intervals are not part of the stored prediction table.
3. **Future outcome uncertainty:** The workflow does not sample binary Bernoulli realizations  $Y_i^{(s)} \sim \text{Bernoulli}(p_i^{(s)})$ ; predictions represent expected conditional probabilities, not simulated futures.
4. **Performance metric uncertainty:** Uncertainty in ROC-AUC, PR-AUC, and metric differences is evaluated via spatial non-parametric bootstrap resampling of the test cells, producing 95% bootstrap confidence intervals.

### 3.4 Decomposition of linear predictor variance

To quantify the shifting relative influence of predictor groups across along-belt sectors, the out-of-fold linear predictor is decomposed additively for each test cell  $i$  in fold  $k$ :

$$\eta_i = c_{\text{intercept},i} + c_{\text{fault},i} + c_{\text{lith},i} + c_{\text{grav},i} + c_{\text{rocks},i}, \quad (11)$$

where each component represents the posterior expectation across retained MCMC draws:

$$c_{\text{fault},i} = \mathbb{E}_{\text{post}} [\beta_{f,\text{lin},d(i)} z_{f,i} + \beta_{f,\text{sq},d(i)} z_{f,i}^2], \quad (12)$$

$$c_{\text{lith},i} = \mathbb{E}_{\text{post}} [\beta_{l,\text{lin},d(i)} z_{l,i} + \beta_{l,\text{sq},d(i)} z_{l,i}^2], \quad (13)$$

$$c_{\text{grav},i} = \mathbb{E}_{\text{post}} [\beta_g z_{g,i}], \quad (14)$$

$$c_{\text{rocks},i} = \mathbb{E}_{\text{post}} [\mathbf{x}_{\text{rock},i}^T \boldsymbol{\beta}_{\text{rock}}]. \quad (15)$$

For each predictor component  $m \in \{\text{fault, lith, grav, rocks}\}$ , its dispersion across the  $N_{\text{test}}$  held-out cells of fold  $k$  is measured by its sample standard deviation:

$$\text{SD}_m = \sqrt{\frac{1}{N_{\text{test}}} \sum_{i \in \mathcal{D}_{\text{test}}^{(k)}} (c_{m,i} - \bar{c}_m)^2}. \quad (16)$$

The relative contribution share of component  $m$  in fold  $k$  is defined as the proportion of total linear-predictor standard deviation attributable to that component:

$$\text{Share}_m = \frac{\text{SD}_m}{\text{SD}_{\text{fault}} + \text{SD}_{\text{lith}} + \text{SD}_{\text{grav}} + \text{SD}_{\text{rocks}}} \times 100\%. \quad (17)$$

These values represent contribution shares to the dispersion of the out-of-fold linear predictor across test cells. They do not represent causal geological effects, Shapley values, or proportions of outcome variance explained ( $R^2$ ).

### 3.5 Distance-representation sensitivity design and robustness criteria

To test whether distance-response inference depends on mathematical representation, four models form a  $2 \times 2$  factorial grid:

Table 1: Distance-representation sensitivity design ( $2 \times 2$  grid).

|                                            | Quadratic     | Linear  |
|--------------------------------------------|---------------|---------|
| Raw distance ( $z$ )                       | Model A (V11) | Model C |
| Log distance ( $\log(1 + x_{\text{km}})$ ) | Model B       | Model D |

Models A and C use standardized raw distance; Models B and D use standardized  $\log(1 + x_{\text{km}})$ . The compact M5 baseline is a non-hierarchical logistic model containing only standardized Bouguer gravity and lithological-contact distance with its quadratic term.

Across the four spatial folds, six Daly domains, and two distance predictors (fault and lithological contact), there are  $4 \times 6 \times 2 = 48$  evaluable comparison cases. In each case, Model A (raw-quadratic) is compared against Model B (log-quadratic) using three pre-registered criteria:

1. **Criterion 1 (Curvature probability agreement):** The absolute difference in posterior probability of positive quadratic curvature between representations must be less than 0.15:

$$|P(\beta_{\text{sq}} > 0 \mid \text{Model A}) - P(\beta_{\text{sq}} > 0 \mid \text{Model B})| < 0.15. \quad (18)$$

2. **Criterion 2 ( $D^*$  median agreement):** The relative difference between posterior median stationary points  $D_A^*$  and  $D_B^*$  (in physical kilometers) must not exceed 25%:

$$\frac{|D_A^* - D_B^*|}{\max(|D_A^*|, |D_B^*|)} \leq 0.25. \quad (19)$$

3. **Criterion 3 (Data support concentration):** Both median stationary points must fall within their respective training data empirical support ranges:

$$\text{support\_min}_A \leq D_A^* \leq \text{support\_max}_A \quad \text{and} \quad \text{support\_min}_B \leq D_B^* \leq \text{support\_max}_B. \quad (20)$$

A relationship is classified as representationally robust if and only if it satisfies all three criteria simultaneously.

## 4 Results

### 4.1 Multivariate posterior structure

The fitted V11 model jointly estimates fault distance, lithological-contact distance, Bouguer gravity, and host lithology across domains and folds. In NRB\_3a, lithological-contact distance consistently displays negative linear coefficients (medians  $-0.849$  to  $-0.451$ ) and positive quadratic curvature (medians  $+0.293$  to  $+0.334$ ), with posterior probabilities of positive curvature  $P(\beta_{l,\text{sq}} > 0) \geq 0.993$  across all four folds (Figure 4). In NRB\_3b, lithological curvature remains positive but exhibits wider posterior uncertainty.




Figure 3: Fold-specific posterior response curves for fault distance across Daly domains under the V11 model. Shaded bands represent posterior credible intervals, illustrating substantial variation across along-belt sectors.

Fault-distance effects vary more dramatically (Figure 3). In NRB\_3a, linear medians range from  $-2.681$  to  $-0.703$ , while quadratic medians are negative in Fold 2 and positive in Folds 1, 3, and 4. In NRB\_3b, fault linear medians shift to positive values ( $+0.181$  to  $+1.548$ ), demonstrating that the directional association with fault distance is not spatially uniform across the belt.




Figure 4: Fold-specific posterior response curves for lithological-contact distance across Daly domains under the V11 model. NRB\_3a exhibits recurring positive quadratic curvature, whereas NRB\_3b shows broader dispersion.

## 4.2 Spatial heterogeneity in predictor contributions

Variance decomposition of the out-of-fold linear predictor components reveals shifting geological and geophysical controls across along-belt sectors (Table 2 and Figure 5):

Table 2: Variance shares of decomposed out-of-fold linear-predictor components.

| Fold | Fault distance | Lithology distance | Bouguer gravity | Host lithology |
|------|----------------|--------------------|-----------------|----------------|
| 1    | 86.5%          | 11.1%              | 2.0%            | 0.4%           |
| 2    | 50.3%          | 23.9%              | 5.6%            | 20.3%          |
| 3    | 33.1%          | 28.2%              | 15.3%           | 23.4%          |
| 4    | 34.9%          | 17.7%              | 31.1%           | 16.3%          |

In Fold 1 (northern sector), fault distance accounts for 86.5% of the linear predictor variance. Moving southward, fault dominance drops to 50.3% in Fold 2, 33.1% in Fold 3, and 34.9% in Fold 4. Concurrently, Bouguer gravity’s contribution increases from 2.0% in Fold 1 to 31.1% in Fold 4, where gravity is nearly co-dominant with fault distance. Host lithology and contact distance provide their largest contributions in Folds 2 and 3.

## 4.3 Distance-response sensitivity and diagnostic failure of $D^*$

The  $2 \times 2$  representation grid demonstrated that mean spatial discrimination is robust to functional form: mean ROC-AUC values are 0.686 for Model A, 0.673 for Model B, 0.671 for Model C, and 0.683 for Model D. Mean PR-AUC values are 0.107, 0.123, 0.130, and 0.145, with the log-linear specification (Model D) achieving the highest PR-AUC.




Figure 5: Variance decomposition of the decomposed out-of-fold linear predictor across the four spatial folds. Fault proximity dominates in the north (Fold 1), while lithology and Bouguer gravity contribute progressively larger shares in central and southern sectors.

While proximity associations persist, polynomial curvature and the derived stationary point  $D^* = -\beta_{\text{lin}}/(2\beta_{\text{sq}})$  are representation-dependent. Pre-registered robustness criteria across the 48 evaluation cases passed in 13 of 48 cases for curvature sign agreement (Criterion 1), 0 of 48 cases for stationary-point agreement (Criterion 2), and 15 of 48 cases for posterior support concentration (Criterion 3). All three criteria were satisfied simultaneously in **0 of 48 cases**.

At the population level, posterior curvature distributions overlap zero:  $P(\mu_{f,\text{sq}} > 0) = 0.384$  and  $P(\mu_{l,\text{sq}} > 0) = 0.743$ . The 95% credible intervals for population  $D^*$  span  $[-279.36, 366.38]$  km for fault distance and  $[-161.69, 152.34]$  km for lithological contact distance. These unphysically wide intervals include negative distances and values far exceeding the basin width, demonstrating that population-level  $D^*$  is statistically unidentifiable and cannot be interpreted as a physical trap distance.

#### 4.4 Spatial predictive transferability and prospectivity mapping

The out-of-fold prospectivity predictions produce the regional prospectivity map shown in Figure 6. Across the four spatial folds, V11 achieves fold-level ROC-AUC values of 0.817, 0.688, 0.676, and 0.561, with PR-AUC values of 0.048, 0.103, 0.133, and 0.145. Pooled across all 1,872 cells, V11 achieves ROC-AUC of 0.689 [95% bootstrap CI: 0.647, 0.727], compared with 0.525 [0.470, 0.581] for the M5 baseline, yielding a pooled gain of +0.164 [+0.084, +0.236].

Pooled performance masks geographic divergence: V11 strongly outperforms M5 in Fold 3 ( $\Delta\text{AUC} = +0.126$ ), whereas M5 outperforms V11 in Fold 2 ( $\Delta\text{AUC} = -0.283$ ), while Folds 1 and 4 yield overlapping confidence intervals. This geographic discrepancy underscores that pooled belt-wide metrics conceal localized predictive failure.

In secondary domain stratification, V11 achieves ROC-AUC of 0.6617 in NRB.3a and 0.5838 in NRB.3b. In single-class domains (CRZ, NKB, SRB, MMSB), ROC-AUC is undefined due to the absence of documented occurrences.




Figure 6: Out-of-fold Bayesian conditional prospectivity map ( $\bar{p}_i$ ) across the Central African Copperbelt under the V11 model. Continuous predicted probabilities reflect spatial variation in prospective potential under along-belt hold-out evaluation.

## 5 Interpretation and representation sensitivity

### 5.1 Conditional multivariate inference and predictor association

The multivariate results establish that geological and geophysical predictors contain joint predictive signal in the evaluated Copperbelt modelling frame. Fault distance, lithological-contact distance, Bouguer gravity, and host lithology contribute jointly to the fitted out-of-fold predictor. The hierarchical structure permits domain-varying intercepts and distance effects while sharing information across Daly domains. This conditional structure means that individual coefficients reflect predictive associations holding other variables constant, rather than independent marginal relationships.

### 5.2 Spatial non-stationarity across along-belt sectors

The strongest empirical finding is the marked spatial variation in predictor contribution shares along the belt. Structural proximity dominates the northern sector (Fold 1), whereas stratigraphic contact proximity, host-rock units, and regional Bouguer gravity provide substantially larger contribution shares in central and southern sectors (Folds 2–4). In the southern fold, Bouguer gravity accounts for 31.1% of linear predictor dispersion, co-dominating with fault distance. This pattern demonstrates that the empirical predictive expression of geological controls is spatially heterogeneous across the belt, consistent with regional variations in basin architecture documented in the Copperbelt geological literature (Cailteux et al., 2005; Selley et al., 2005; Hitzman et al., 2012; Ndonfack et al., 2025).

### 5.3 Representation sensitivity and diagnostic role of $D^*$

Proximity associations persist across mathematical transformations, but polynomial curvature and the derived stationary point  $D^*$  are sensitive to functional form. The zero-pass rate across the 48 robustness cases demonstrates that  $D^*$  is a model-dependent algebraic property of quadratic curves rather than an invariant physical optimum. Consequently,  $D^*$  should be treated as an audit diagnostic for polynomial parameterizations, and prospectivity mapping should avoid translating quadratic extrema into universal exploration drilling targets.

## 6 Discussion, limitations, and exploration implications

### 6.1 What the model establishes: empirical association versus causal ore genesis

The results demonstrate that geological proxies contain predictive signal for sediment-hosted copper prospectivity, but their relationships are spatially non-stationary. The shifting variance shares reflect changing empirical predictor importance across sectors. We caution against conflating statistical variance decomposition with causal ore-forming mechanisms. The increased contribution of Bouguer gravity in Fold 4 does not imply that gravity caused mineralization; rather, regional density contrasts provide a stronger conditional predictor of deposit location in the southern sector where fault proximity is less discriminative.

## 6.2 Regional transferability and Daly’s tectonic framework

Spatially separated prediction demonstrates that pooled belt-wide skill can conceal regional performance divergence. V11 outperforms the baseline in Fold 3 but underperforms it in Fold 2. This divergence demonstrates that prospectivity models cannot be assumed to transfer uniformly across a mineralized belt, reinforcing the necessity of spatially separated validation strategies (Roberts et al., 2017; Valavi et al., 2019).

Daly’s tectonic domain classification (Daly et al., 2025) provides a principled geological stratification for hierarchical partial pooling. While the model accommodates domain-level parameter differences, the analysis estimates empirical associations within the compiled data and does not constitute a formal proof of Daly’s tectonic hypothesis in its entirety.

## 6.3 Model parsimony and distance representations

The comparable discrimination achieved by linear and log-linear specifications (Models C and D) relative to quadratic models indicates that quadratic distance terms are not universally required to achieve predictive skill. Linear and logarithmic distance representations offer greater parsimony and avoid introducing unidentifiable polynomial stationary points, making them preferable baselines for regional prospectivity screening.

## 6.4 Methodological and data limitations

Several methodological limitations must be acknowledged:

1. **Sample size and domain imbalance:** Only 138 deposit-present cells are available, with 107 concentrated in NRB\_3a and 31 in NRB\_3b. Four Daly domains contain zero positive observations in this modelling frame, precluding within-domain discrimination evaluation and restricting parameter inference in those domains to hierarchical shrinkage.
2. **Residual spatial autocorrelation:** Although contiguous spatial blocking mitigates optimistic cross-validation bias, spatial autocorrelation may persist across block boundaries (Roberts et al., 2017).
3. **Input layer uncertainty:** Regional geological maps, structural compilations, and gravity grids contain measurement, interpolation, and cartographic errors that propagate into distance calculations.
4. **Surface proxies vs. 3D geometry:** Near-surface distance features do not capture three-dimensional fault dip, subsurface stratigraphy, or depth to basement.
5. **Discrete blocking approximation:** Four discrete spatial folds approximate what is fundamentally a continuous along-belt metallogenic gradient.
6. **Reporting scope of uncertainty:** Stored out-of-fold prospectivity predictions represent expected conditional probabilities  $\bar{p}_i$ ; cell-level credible intervals and simulated Bernoulli outcome distributions ( $Y_i^{(s)} \sim \text{Bernoulli}(p_i^{(s)})$ ) were not part of the frozen reporting tables.

## 7 Conclusions

An uncertainty-aware hierarchical Bayesian framework enables rigorous evaluation of mineral prospectivity mapping under along-belt spatial separation. In the Central African Copperbelt, predictor associations are geographically heterogeneous: structural proximity dominates predictive discrimination in northern sectors, whereas stratigraphic contacts and regional gravity play larger roles in central and southern sectors. Distance-response associations persist across functional representations, but quadratic curvature and derived  $D^*$  stationary points are representation-dependent and unidentifiable at the population scale. Prospective models should be evaluated under spatially separated cross-validation and interpreted as conditional predictive representations rather than invariant physical laws.

## 8 Reproducibility, data, and code availability

The modeling workflow is implemented in Python using PyMC (Salvatier et al., 2016), ArviZ, and scikit-learn. The core spatial cross-validation script is `src/v11_spatial_stability_final.py`; the distance sensitivity framework is in `src/run_distance_sensitivity.py`; and the population reconstruction pipeline is in `src/run_v11_population_reconstruction.py`. MCMC sampling used the No-U-Turn Sampler (Hoffman and Gelman, 2014) with 2 chains  $\times$  1,500 retained iterations per fold following 1,000 tuning steps (target acceptance rate 0.90, random seed 42). Source spatial data layers and frozen posterior traces are archived within the project repository.

## References

- Graeme F. Bonham-Carter. *Geographic Information Systems for Geoscientists: Modelling with GIS*. Pergamon, Oxford, 1994.
- Graeme F. Bonham-Carter, Frederik P. Agterberg, and David F. Wright. Weights of evidence modelling: A new approach to mapping mineral potential. In Frederik P. Agterberg and Graeme F. Bonham-Carter, editors, *Statistical Applications in the Earth Sciences*, Paper 89-9, pages 171–183. Geological Survey of Canada, 1989. doi: 10.4095/128059.
- Alexander Brenning. Spatial cross-validation and bootstrap for the assessment of prediction rules in remote sensing: The R package `sperrorest`. *Proceedings of the IEEE International Geoscience and Remote Sensing Symposium (IGARSS)*, pages 5372–5375, 2012. doi: 10.1109/IGARSS.2012.6352393.
- Jacques L. H. Cailteux, Abderrahmane B. Kampunzu, Catherine Lerouge, Alexis K. Kaputo, and Jean P. Milesi. Genesis of sediment-hosted stratiform copper–cobalt deposits, Central African Copperbelt. *Journal of African Earth Sciences*, 42(1–5):134–158, 2005. doi: 10.1016/j.jafrearsci.2005.08.001.
- Emmanuel John M. Carranza. *Geochemical Anomaly and Mineral Prospectivity Mapping in GIS*. Elsevier, Amsterdam, 1st edition, 2009. doi: 10.1016/S0169-7161(09)00206-1.
- Emmanuel John M. Carranza and Alice G. Laborte. Random forest predictive modeling of mineral prospectivity with small number of prospects and data with missing values in Abra (Philippines). *Computers and Geosciences*, 74:60–70, 2015. doi: 10.1016/j.cageo.2014.10.004.

- M. C. Daly, M. Zimba, M. Purkiss, and F. Chibesakunda. The inter-cratonic Neoproterozoic Katangan basin of central Africa: Rift basin formation, orogenic inversion and potential fluid pathways. *Tectonics*, 44, 2025. doi: 10.1029/2025TC008955.
- Alistair Ford. Practical implementation of random forest-based mineral potential mapping for porphyry Cu–Au mineralization in the Eastern Lachlan Orogen, NSW, Australia. *Natural Resources Research*, 29:267–283, 2019. doi: 10.1007/s11053-019-09598-y.
- A. Stewart Fotheringham, Chris Brunsdon, and Martin Charlton. *Geographically Weighted Regression: The Analysis of Spatially Varying Relationships*. Wiley, Chichester, 2002.
- Alan E. Gelfand, Hyon-Jung Kim, C. F. Sirmans, and Sudipto Banerjee. Spatial modeling with spatially varying coefficient processes. *Journal of the American Statistical Association*, 98(462): 387–396, 2003. doi: 10.1198/016214503000170.
- Andrew Gelman and Jennifer Hill. *Data Analysis Using Regression and Multilevel/Hierarchical Models*. Cambridge University Press, Cambridge, 2007. doi: 10.1017/CBO9780511790942.
- Murray W. Hitzman, David Broughton, David Selley, Jon Woodhead, David Wood, and Stuart Bull. The Central African Copperbelt: Diverse stratigraphic, structural, and temporal settings in the world’s largest sedimentary copper district. In Jeffrey W. Hedenquist, Mark Harris, and Francisco Camus, editors, *Geology and Genesis of Major Copper Deposits and Districts of the World: A Tribute to Richard H. Sillitoe*, volume 16 of *Special Publication*, pages 487–514. Society of Economic Geologists, 2012. doi: 10.5382/SP.16.19.
- Matthew D. Hoffman and Andrew Gelman. The No-U-Turn sampler: Adaptively setting path lengths in Hamiltonian Monte Carlo. *Journal of Machine Learning Research*, 15:1593–1623, 2014.
- R. M. Key, B. De Waele, and A. K. Liyungu. A multi-element baseline geochemical database from the western extension of the Central Africa Copperbelt in northwestern Zambia. *Applied Earth Science*, 113(3):B205–B212, 2004. doi: 10.1179/037178404225005437.
- Vladimir A. Lisitsin, Alok Porwal, and T. Campbell McCuaig. Probabilistic fuzzy logic modeling: Quantifying uncertainty of mineral prospectivity models using Monte Carlo simulations. *Mathematical Geosciences*, 46(6):747–769, 2014. doi: 10.1007/s11004-014-9504-5.
- Jorge M. Lobo, Alberto Jiménez-Valverde, and Raimundo Real. AUC: A misleading measure of the performance of predictive distribution models. *Global Ecology and Biogeography*, 17(2):145–151, 2008. doi: 10.1111/j.1466-8238.2007.00358.x.
- Tshanga Matthieu Matthieu, Lindani Ncube, and Kgabo Humphrey Thamaga. Geological mapping of copper deposits in the Democratic Republic of Congo through remote sensing data and machine learning. *Frontiers in Remote Sensing*, 6:1678991, 2026. doi: 10.3389/frsen.2025.1678991.
- Kevin Igor Azeuda Ndonfack, Zhiming Yang, Jinlin Zhang, Scott A. Whattam, and Yuling Xie. Geology, geochemistry, and exploration of the Central African Copperbelt: A review. *International Geology Review*, 67(8), 2025. doi: 10.1080/00206814.2024.2426200.
- Mohammad Parsa and Emmanuel John M. Carranza. Modulating the impacts of stochastic uncertainties linked to deposit locations in data-driven predictive mapping of mineral prospectivity. *Natural Resources Research*, 30(5):3081–3097, 2021. doi: 10.1007/s11053-021-09891-9.

- Alok Porwal, Emmanuel John M. Carranza, and Martin Hale. Bayesian network classifiers for mineral potential mapping. *Computers & Geosciences*, 32(1):1–16, 2006. doi: 10.1016/j.cageo.2005.03.018.
- David R. Roberts, Volker Bahn, Simone Ciuti, Mark S. Boyce, Jane Elith, Gurutzeta Guillera-Arroita, Severin Hauenstein, José J. Lahoz-Monfort, Boris Schröder, Wilfried Thuiller, David I. Warton, Brendan A. Wintle, Florian Hartig, and Carsten F. Dormann. Cross-validation strategies for data with temporal, spatial, hierarchical, or phylogenetic structure. *Ecography*, 40(8):913–929, 2017. doi: 10.1111/ecog.02881.
- John Salvatier, Thomas V. Wiecki, and Christopher Fonnesbeck. Probabilistic programming in Python using PyMC3. *PeerJ Computer Science*, 2:e55, 2016. doi: 10.7717/peerj-cs.55.
- Roger Sayre, Patrick Comer, Jon Hak, Carmen Josse, Jennifer Bow, Harumi Warner, Mahamane Larwanou, Ensermu Kelbessa, Tamrat Bekele, Henry Kehl, Rene Colditz, Joyce Fry, and William Kramber. *A New Map of Global Ecological Land Units — An Ecologically Meaningful Classification of the Earth’s Surface*. Association of American Geographers and U.S. Geological Survey, Washington, DC, 2014.
- David Selley, David Broughton, Robert J. Scott, Murray W. Hitzman, Stuart W. Bull, Ross R. Large, Peter J. McGoldrick, Mawson Croaker, Noel Pollington, and Fernando Barra. A new look at the geology of the Zambian Copperbelt. In Jeffrey W. Hedenquist, John F. H. Thompson, Richard J. Goldfarb, and Jeremy P. Richards, editors, *Economic Geology 100th Anniversary Volume*, pages 965–1000. Society of Economic Geologists, 2005. doi: 10.5382/AV100.29.
- Roozbeh Valavi, Jane Elith, José J. Lahoz-Monfort, and Gurutzeta Guillera-Arroita. blockCV: An R package for generating spatially or environmentally separated folds for k-fold cross-validation of species distribution models. *Methods in Ecology and Evolution*, 10(2):225–232, 2019. doi: 10.1111/2041-210X.13107.
- Mahyar Yousefi, Emmanuel John M. Carranza, Oliver P. Kreuzer, Vesa Nykänen, Jon M. A. Hronsky, and Mark J. Mihalasky. Data analysis methods for prospectivity modelling as applied to mineral exploration targeting: State-of-the-art and outlook. *Journal of Geochemical Exploration*, 229:106839, 2021. doi: 10.1016/j.gexplo.2021.106839.
- Renguang Zuo, Oliver P. Kreuzer, Jian Wang, Yihui Xiong, Zhenjie Zhang, and Ziyue Wang. Uncertainties in GIS-based mineral prospectivity mapping: Key types, potential impacts and possible solutions. *Natural Resources Research*, 30(5):3059–3079, 2021. doi: 10.1007/s11053-021-09871-z.